

PsychiatryChat

Game Design and Technical Design Document

Document purpose: Provide Codex with the design requirements and broad technical architecture needed to build a functional prototype of PsychiatryChat.

Product type: Single-player, browser-based educational narrative game

Primary audience: Medical students and psychiatry residents

Platform: Desktop and mobile web browser

Initial delivery format: HTML, CSS, and JavaScript web application that can run locally without a server

Visual direction: A warm, atmospheric, conversation-centered visual-novel presentation broadly reminiscent of games such as *Coffee Talk*, while using entirely original layouts, artwork, characters, writing, and interface elements.

1. Project Summary

PsychiatryChat is a longitudinal psychiatry communication game in which the player assumes the role of a psychiatry trainee clinician.

At the beginning of the game, the player chooses one of three patients. The player then follows that patient over a six-week period through three clinical encounters:

1. Initial encounter
2. Midpoint follow-up
3. Final encounter

Between encounters, predetermined events occur in the patient's life. The external events themselves are fixed, but the patient's interpretation of, reaction to, and behavior following those events are influenced by the clinician's earlier dialogue choices.

The central teaching premise is that clinical language is not merely descriptive. What a clinician says—and how they say it—can influence:

- What a patient feels safe disclosing
- Whether a patient feels understood or judged
- Treatment engagement
- Medication adherence
- Help-seeking behavior

- Family or social communication
- The patient's interpretation of symptoms
- The therapeutic alliance
- How the patient responds to later stressors
- Whether a rupture is repaired or compounded

The game should not suggest that a clinician single-handedly controls a patient's life or psychiatric outcome. Patient outcomes are multifactorial. The game instead demonstrates that therapeutic communication can meaningfully alter one part of a larger causal system.

2. High-Level Design Goals

2.1 Primary goals

PsychiatryChat should:

- Demonstrate the longitudinal consequences of clinician communication.
- Allow players to see how apparently small wording differences can influence later disclosure, trust, agency, and treatment engagement.
- Present psychiatry as a relationship-based clinical discipline rather than a sequence of diagnostic multiple-choice questions.
- Encourage reflection rather than simple answer memorization.
- Include ambiguity, imperfect choices, and opportunities to repair prior mistakes.
- Reward curiosity, validation, directness, collaboration, and appropriate safety assessment without reducing every conversation to a visible score.
- Make replaying the same case educational by revealing different patient trajectories and endings.

2.2 Secondary goals

The game should also teach:

- Therapeutic interviewing
- Shared decision-making
- Risk and safety assessment
- Trauma-informed communication
- Appropriate use of silence
- Validation without falsely endorsing distorted beliefs
- Boundary setting
- Collaborative treatment planning
- Recognition and repair of therapeutic rupture
- Patient autonomy
- The difference between empathy and avoidance
- The difference between reassurance and invalidation
- How diagnostic and treatment language may be received by a patient

2.3 Non-goals

The initial version should not attempt to be:

- A comprehensive psychiatry curriculum
- A full diagnostic simulator
- A prescribing simulator
- A replacement for standardized patients or supervised clinical experience
- A chatbot powered by generative AI
- An open-ended natural-language conversation engine
- A randomized roguelike
- A system in which every line of dialogue has a clearly correct or incorrect answer

The first version should use carefully authored dialogue choices and deterministic branching.

3. Core Player Experience

3.1 Opening flow

The game opens with:

1. Title screen
2. Brief explanation of the six-week structure
3. Educational disclaimer
4. Patient selection screen
5. Selected patient's introductory profile
6. Encounter 1

The player chooses one of three patient cards. Each card should provide enough information to distinguish the cases without revealing the central conflict or possible endings.

Each patient card should include:

- First name
- Age
- Pronouns
- Clinical setting
- Presenting concern
- One-sentence personal context
- Approximate play time
- Major teaching themes

Example:

Maya, 26

Outpatient psychiatry

Increasing anxiety and missed work

Themes: validation, avoidance, medication concerns, therapeutic alliance

The three initial patients should be new characters. The attached cases may provide loose inspiration regarding depression, medication planning, mania, safety, and clinical communication, but the game should not directly reproduce their dialogue or exact structure.

4. Six-Week Timeline

The game covers six fictional weeks but should take approximately 25–40 minutes to complete for one patient.

Week 1: Encounter 1

Purpose:

- Establish the therapeutic relationship.
- Introduce the patient's current symptoms and life context.
- Allow the player's communication style to influence what is disclosed.
- Establish the first major trajectory branch.

Encounter 1 should contain approximately:

- 20–30 dialogue exchanges
- 6–9 meaningful player decisions
- Several non-branching conversational beats
- One hidden branch resolution at the end

The player should not be told, "You selected Branch A." The transition should be represented through the patient's behavior, internal state, or subsequent life-event response.

Weeks 2–3: First inter-encounter period

A fixed external event occurs.

Examples:

- Conflict with a supervisor
- Missed medication doses
- Argument with a family member
- Unexpected contact from a former partner
- Academic setback
- Distressing physical symptom

- Loss of housing stability
- Social invitation
- Increased work responsibility

The event occurs in every playthrough of that patient's case.

The patient's response changes based on the Encounter 1 trajectory.

For example:

Same event: The patient receives critical feedback at work.

Trajectory A response: The patient recalls the clinician's collaborative framing, asks for clarification, and contacts the clinic when symptoms worsen.

Trajectory B response: The patient interprets the criticism as confirmation of personal failure, conceals worsening symptoms, and stops responding to clinic messages.

The game presents this period as a short visual montage rather than a full clinical encounter.

Week 3: Encounter 2

Purpose:

- Reveal the consequences of the initial interaction.
- Allow the player to reinforce, redirect, or repair the therapeutic relationship.
- Introduce new clinical information based on what the patient is now willing to disclose.
- Establish the second major trajectory branch.

Encounter 2 should contain approximately:

- 20–30 dialogue exchanges
- 6–9 meaningful decisions
- Different opening dialogue depending on the Encounter 1 branch
- One hidden branch resolution

Encounter 2 must allow repair. A weak first encounter should not permanently force the worst ending.

Likewise, a strong first encounter should not guarantee the best ending. A player can mishandle new information, become overly directive, minimize a concern, or fail to respond to a rupture.

Weeks 4–5: Second inter-encounter period

A second fixed life event occurs.

Its outcome presentation depends on both prior trajectory resolutions.

At this stage there are four authored versions of the patient's response:

- $A \rightarrow A$
- $A \rightarrow B$
- $B \rightarrow A$
- $B \rightarrow B$

These four combinations lead to four distinct final encounter states and four eventual endings.

Week 6: Encounter 3

Purpose:

- Show the cumulative result of the therapeutic process.
- Allow the player to speak with the patient after the major outcome has already begun to emerge.
- Provide emotional and educational closure.
- Demonstrate how the clinician's prior words remained relevant outside the clinic.
- Deliver a case-specific reflection.

Encounter 3 may include dialogue choices, but these choices should not create another major life-outcome branch in the initial version. The four canonical endings should already be determined by the first two encounter resolutions.

Final-encounter choices may affect:


- The tone of the closing conversation
- Whether the player recognizes what happened
- Whether the patient feels understood at the end
- The quality of the player's final reflection
- Which teaching notes are displayed
- A small epilogue variation

They should not create more than four principal endings per patient.

5. Four-Ending Branch Structure

Each patient has exactly two major branch resolutions.

```
Encounter 1
├─ Trajectory A
│   └─ Encounter 2
│       ├─ Trajectory A → Ending AA
│       └─ Trajectory B → Ending AB
│
```



This produces four endings for each patient and twelve principal endings across the three-patient game.

5.1 Meaning of the branches

The branches should be case-specific rather than universally labeled “good” and “bad.”

At the structural level:

Ending AA: Sustained therapeutic momentum

The initial encounter establishes a useful alliance, and the second encounter responds effectively to new stress or information.

This is generally the most constructive outcome, but it should not be unrealistically perfect.

Ending AB: Early alliance followed by rupture or missed opportunity

The player begins well but mishandles a later disclosure, becomes overly directive, misses a safety issue, or fails to adapt when circumstances change.

This teaches that an alliance must be maintained rather than earned once.

Ending BA: Delayed repair and re-engagement

The first encounter is incomplete, invalidating, or overly clinical, but the player recognizes the rupture and responds differently during Encounter 2.

This should be a meaningful and often emotionally satisfying ending. It reinforces that repair is a central therapeutic skill.

Ending BB: Compounding disengagement or harm

The initial communication problem is not repaired and is reinforced during the second encounter.

The resulting outcome may involve disengagement, worsening symptoms, delayed help-seeking, poor adherence, escalating conflict, or emergency care.

The ending should remain plausible and respectful. It should not exist merely to punish the player.

6. Interpreting the “50/50” Design

The 50/50 structure refers to two possible authored trajectories after Encounter 1 and two possible trajectories after Encounter 2.

It should not mean that the game randomly determines whether a patient improves.

Branching must be determined by the player’s cumulative dialogue choices.

Each encounter should contain multiple choices that influence the branch resolution. No single ordinary dialogue option should visibly function as a giant “good ending versus bad ending” switch.

A safety-critical choice may have disproportionate weight, but the overall branch should generally reflect a pattern of interaction.

Example:

Encounter 1 hidden dimensions:

- Trust
- Felt validation
- Patient agency
- Disclosure
- Safety clarity
- Treatment collaboration

The encounter resolver evaluates the relevant case-specific dimensions at the end of the encounter.

The player never sees numerical values during the initial playthrough.

7. Dialogue Mechanics

7.1 Basic conversation loop

Each scene should follow this pattern:

1. Patient speaks.
2. Patient portrait or expression changes.
3. Player receives two to four possible responses.
4. Player selects a response.
5. Patient reacts.
6. Hidden state changes.
7. Conversation continues.

7.2 Types of dialogue choices

Choices should include a mixture of:

- Open-ended questions
- Closed questions
- Reflective statements
- Validation
- Clarification
- Direct safety questions
- Premature reassurance
- Interpretation
- Advice
- Collaborative planning
- Directive planning
- Boundary setting
- Silence or pause
- Deflection
- Medicalized language
- Everyday language
- Appropriate challenge
- Inappropriate confrontation

Options should be credible enough that a learner might genuinely consider more than one.

Avoid choice sets in which one response is obviously compassionate and the others are absurdly hostile.

Weak options should generally represent realistic trainee mistakes:

- Moving too quickly
- Trying to fix the problem before understanding it
- Softening a necessary safety question
- Overusing diagnostic language
- Asking several questions at once
- Offering reassurance that minimizes the patient's experience
- Prioritizing checklist completion over the relationship
- Becoming paternalistic
- Avoiding an uncomfortable topic
- Confusing validation with agreement
- Failing to acknowledge a rupture

7.3 Response timing

Do not display immediate "Correct" or "Incorrect" labels during the primary playthrough.

The patient's response should be the first form of feedback.

Example:

Player: "I'm sure your manager didn't mean it that way."

Patient: "Maybe. I guess it sounds stupid when I explain it."

This reaction communicates that the reassurance closed down disclosure without interrupting the scene with a lecture.

7.4 Delayed educational feedback

Formal teaching should occur:

- At the end of each encounter
- After inter-encounter consequences are shown
- At the final case debrief
- In an optional educator/replay mode

This preserves narrative immersion while still making the educational content explicit.

8. Hidden State Model

Each case should define its own state variables. The engine should not assume that all patients use the same clinical dimensions.

A shared base state may include:

```
{
  selectedPatientId: "",
  currentWeek: 1,
  currentEncounter: 1,
  currentSceneId: "",
  encounterHistory: [],
  choiceHistory: [],
  flags: {},
  dimensions: {},
  firstBranch: null,
  secondBranch: null,
  endingId: null,
  completed: false
}
```

Possible hidden dimensions include:

```

{
  trust: 0,
  disclosure: 0,
  agency: 0,
  feltValidation: 0,
  shame: 0,
  defensiveness: 0,
  treatmentBuyIn: 0,
  safetyClarity: 0,
  hope: 0,
  dependency: 0,
  avoidance: 0,
  rupture: 0,
  repair: 0
}

```

Not every case should use every dimension.

Each dialogue choice can change one or more dimensions:

```

{
  id: "maya_e1_04_reflect",
  text: "It sounds like you have been trying to hold this together by
yourself.",
  effects: {
    trust: 1,
    feltValidation: 2,
    disclosure: 1
  },
  addFlags: ["recognized_isolation"],
  teachingTags: ["reflection", "validation"],
  nextSceneId: "maya_e1_05"
}

```

A more directive option might have mixed effects:

```

{
  id: "maya_e1_04_plan",
  text: "Let's make a plan for getting you back to work this week.",
  effects: {
    hope: 1,
    agency: -1,
    disclosure: -1
  },
}

```

```
addFlags: ["premature_problem_solving"],
teachingTags: ["advice", "timing"],
nextSceneId: "maya_e1_05_plan_response"
}
```

Mixed effects are important. An option can be well-intentioned and partially useful while still creating a relational cost.

9. Encounter Resolution

Each encounter should have a case-specific resolution function.

Example:

```
function resolveEncounterOne(state, patientConfig) {
  const allianceScore =
    state.dimensions.trust +
    state.dimensions.feltValidation +
    state.dimensions.disclosure +
    state.dimensions.agency -
    state.dimensions.shame -
    state.dimensions.rupture;

  const requiredSafetyFlag =
    !patientConfig.requiresSafetyAssessment ||
    state.flags.directSafetyAssessmentCompleted;

  return allianceScore >= patientConfig.encounterOneThreshold &&
    requiredSafetyFlag
    ? "A"
    : "B";
}
```

Codex may refine the exact implementation, but the resolver should:

- Be deterministic
- Use several choices rather than one arbitrary switch
- Support mandatory case-specific flags
- Allow mixed responses
- Be easy for educators to inspect
- Be easy to rebalance without rewriting the engine

For some cases, missing a critical safety assessment may cap the possible branch even when the relational score is otherwise strong.

10. Static Life Events and Dynamic Patient Responses

Each patient should have two fixed inter-encounter events.

The event data structure should separate the objective event from the patient's possible reactions.

```
{
  id: "maya_event_1",
  weekLabel: "Two weeks later",
  eventTitle: "A difficult meeting",
  objectiveEvent: [
    "Maya's supervisor asks to meet about several missed shifts."
  ],
  variants: {
    A: {
      narration: [
        "Maya remembers discussing how avoidance has been narrowing her options.",
        "She attends the meeting and asks whether temporary accommodations are possible."
      ],
      effects: {
        agency: 1,
        treatmentBuyIn: 1
      }
    },
    B: {
      narration: [
        "Maya assumes the meeting will confirm that everyone sees her as unreliable.",
        "She calls out again and does not open the supervisor's follow-up email."
      ],
      effects: {
        avoidance: 2,
        shame: 1
      }
    }
  }
}
```

The second event should support four variants:

```
variants: {  
  AA: {},  
  AB: {},  
  BA: {},  
  BB: {}  
}
```

Presentation

Inter-encounter sequences should feel like short narrative montages.

Possible presentation elements:

- Week marker
- Calendar pages
- Text messages
- Voicemail transcript
- Workplace or home background
- Patient internal narration
- Medication bottle, appointment reminder, or unopened message
- Short third-person narration
- Ambient audio change
- Patient journal excerpt

These events should usually be noninteractive. Their purpose is to show consequences rather than give the player another opportunity to control every part of the patient's life.

A “Continue” button advances the montage.

11. Teaching Integration

11.1 Teaching layer one: Patient reaction

The first and most important teaching mechanism is the patient's response.

The response should show:

- Increased or decreased openness
- Emotional shift
- Changes in body language
- Whether the patient corrects the clinician
- Whether the patient withdraws
- Whether new information is disclosed

- Whether the patient accepts or resists a plan

11.2 Teaching layer two: Encounter reflection

At the end of each encounter, show a brief reflection screen.

Suggested format:

What the patient may have heard

A short interpretation of the interaction from the patient's perspective.

Clinical communication principle

One concise teaching point.

A phrase worth reconsidering

Show one selected player statement and explain its likely effect.

Alternative approach

Provide one example of a stronger response.

Do not review every decision after every encounter. Select two or three pivotal moments.

11.3 Teaching layer three: Longitudinal consequence

The inter-encounter sequence demonstrates how communication may influence later behavior.

This is a core part of the curriculum and should not be treated as decorative story content.

11.4 Teaching layer four: Final debrief

After Encounter 3, display:

- Ending title
- Patient outcome
- Therapeutic relationship outcome
- Major clinical consequence
- Two statements that helped
- Two statements that limited or harmed the encounter
- Missed disclosures or assessments
- Repair opportunities
- A case-specific teaching pearl
- Reminder that the outcome was multifactorial

- Option to view the branch map
- Option to replay

11.5 Educator mode

After completing a case once, unlock an optional educator mode.

Educator mode may show:

- Hidden dialogue tags
- State changes
- Branch requirements
- Pivotal decisions
- Why a particular patient response appeared
- Alternative patient responses
- Learning objectives
- Suggested discussion questions

This information should remain hidden during the first playthrough.

12. Provisional Patient Framework

The exact patient stories can be written later. Codex should build the game so that case content can be replaced without changing the engine.

A provisional three-case structure may be used for testing.

Patient 1: Depression, shame, and disclosure

Possible themes:

- Depressive symptoms
- Work impairment
- Passive death wish or suicidal thinking
- Fear of being a burden
- Direct versus euphemistic safety questions
- Premature reassurance
- Collaborative treatment planning
- The effect of language on later help-seeking

Primary trajectory axis:

- Disclosure and help-seeking versus concealment and withdrawal

Patient 2: Emerging mania or psychosis and autonomy

Possible themes:

- Suspicion
- De-escalation
- Validation without confirming a delusion
- Medication concerns
- Family involvement
- Coercion and autonomy
- Treatment adherence
- Rupture repair

Primary trajectory axis:

- Collaboration and trust versus escalating defensiveness and disengagement

Patient 3: Anxiety, trauma, or recurrent crisis behavior

Possible themes:

- Avoidance
- Shame
- Boundary setting
- Trauma-informed communication
- Clinician frustration
- Dependency versus agency
- Safety planning
- Validation combined with behavioral change

Primary trajectory axis:

- Increased agency and coping versus avoidance, dependency, or repeated crisis use

These are placeholders rather than final case prescriptions.

13. Interface Design


13.1 Main encounter layout

The primary screen should include:

- Atmospheric background
- Patient portrait or half-body character image
- Patient name

- Dialogue text box
- Speaker label
- Two to four response buttons
- Week and encounter indicator
- Menu button
- Optional dialogue history button

Suggested desktop layout:



On mobile, the portrait should occupy the upper portion of the screen and the dialogue interface should remain readable without horizontal scrolling.

13.2 Art placeholders

Until final graphics are provided, use:

- Simple illustrated silhouette portraits
- Expression labels or basic portrait variants
- Gradient or blurred environmental backgrounds
- Clearly named placeholder image files
- Neutral original assets rather than copied game artwork

Expected portrait states may include:

- Neutral
- Guarded
- Anxious
- Irritated
- Tearful
- Relieved
- Thoughtful

- Withdrawn

13.3 Visual tone

The interface should feel:

- Warm
- Quiet
- Intimate
- Reflective
- Slightly nocturnal or atmospheric
- Professional without resembling an electronic medical record

Avoid:

- Cartoonishly exaggerated psychiatric symptoms
 - Horror aesthetics
 - Red flashing “danger” interfaces
 - Gamified patient health bars
 - Visible morality meters
 - “Patient defeated” or “mission failed” language
-

14. Audio

Audio is optional for the first prototype.

The architecture should permit later addition of:

- Low-volume ambient room sound
- Rain, traffic, ventilation, or café-like environmental ambience
- Subtle scene-transition sound
- Distinct musical motifs for each patient
- Audio changes during rupture, disclosure, or reflection

Include:

- Master mute
- Music volume
- Ambient volume
- Reduced sensory mode

Do not require audio to understand any content.

15. Technical Architecture

15.1 General requirements

The prototype should:

- Use standard HTML, CSS, and JavaScript.
- Require no backend.
- Require no account or login.
- Run from a local web server or static hosting service.
- Be deployable to GitHub Pages or similar static hosting.
- Avoid external dependencies when practical.
- Keep case content separate from rendering logic.
- Save progress with `localStorage`.
- Work in current desktop and mobile browsers.
- Support keyboard navigation.
- Be responsive.

A lightweight JavaScript framework is not necessary. Vanilla JavaScript is preferred for the first version unless Codex determines that a small framework materially improves maintainability.

15.2 Suggested file structure

```
psychiatry-chat/  
├─ index.html  
├─ css/  
│   ├─ base.css  
│   ├─ game.css  
│   └─ accessibility.css  
├─ js/  
│   ├─ app.js  
│   ├─ engine.js  
│   ├─ state.js  
│   ├─ renderer.js  
│   ├─ storage.js  
│   ├─ teaching.js  
│   └─ accessibility.js  
├─ data/  
│   ├─ patients.js  
│   ├─ patient-1.js  
│   ├─ patient-2.js  
│   └─ patient-3.js  
├─ assets/  
│   ├─ portraits/  
│   └─ backgrounds/
```

```
|   |— ui/
|   |— audio/
|— README.md
```

The final build may optionally bundle everything into a single self-contained HTML file, but the source version should remain modular.

15.3 Content-driven engine

The engine should render scenes from structured data rather than hard-coded page sequences.

Suggested scene model:

```
{
  id: "maya_e1_scene_04",
  type: "dialogue",
  speaker: "Maya",
  portrait: "maya_guarded",
  background: "clinic_evening",
  text: "I don't know. Maybe I'm just lazy.",
  choices: [
    {
      id: "reflect_shame",
      text: "It sounds like you have been blaming yourself for having a hard
time.",
      nextSceneId: "maya_e1_scene_05a",
      effects: {
        feltValidation: 2,
        disclosure: 1
      },
      flags: ["shame_reflected"],
      teachingTags: ["reflection", "shame"]
    },
    {
      id: "normalize",
      text: "A lot of people struggle with motivation.",
      nextSceneId: "maya_e1_scene_05b",
      effects: {
        feltValidation: 0,
        disclosure: -1
      },
      flags: ["generic_normalization"],
      teachingTags: ["normalization", "missed_emotion"]
    }
  ]
}
```

```
]
}
```

Scene types should include:

- title
- dialogue
- narration
- choice
- lifeEvent
- encounterReflection
- ending
- caseDebrief

15.4 Conditional content

Scenes and choices should support conditions.

```
{
  conditions: [
    { type: "flag", key: "asked_about_safety", equals: true },
    { type: "dimension", key: "trust", operator: ">=", value: 2 }
  ]
}
```

Possible conditional behaviors:

- Show a different patient greeting.
- Unlock a disclosure.
- Change an available response.
- Change portrait expression.
- Modify narration.
- Hide information the player never elicited.
- Reference a prior clinician statement.
- Offer a rupture-repair option.
- Change the final debrief.

15.5 Dialogue memory

The game should remember selected pivotal phrases.

Later patient dialogue may reference them:

"Last time you said we could make the plan together."

or:

“You told me everyone feels like this sometimes.”

This can be implemented by saving specific choice IDs or flags and inserting conditional text later.

The patient should not recall every selection. Only deliberately authored pivotal statements should recur.

16. Saving and Replay

Use `localStorage` to save:

- Selected patient
- Current scene
- Current week
- Current encounter
- Choice history
- Flags
- Hidden dimensions
- Branch results
- Completed endings
- Unlocked educator mode
- Settings

Provide menu options for:

- Continue
- Save and exit
- Restart current patient
- Return to patient selection
- View completed endings
- Change settings

Backtracking

During the initial playthrough, do not allow the player to undo individual dialogue choices.

This reinforces that spoken words cannot simply be taken back.

After the player completes a case, unlock a replay mode that permits:

- Restarting an encounter
- Jumping to previously reached branches
- Viewing the branch map
- Comparing alternative responses

17. Branch Map

After completing a patient once, display a partially revealed branch map.

The map should show:

- The path the player took
- The ending reached
- Unseen branches as silhouettes or question marks
- Previously discovered endings
- Major rupture or repair points after they have been encountered

Do not reveal exact scoring thresholds unless educator mode is enabled.

Example:

```
Encounter 1
├─ Felt understood — Encounter 2 — [Current ending]
│
├─ Withheld concerns — Encounter 2 — ???
│
└─ ???
```

This encourages replay without reducing the game to optimization during the first run.

18. Accessibility

The prototype must include:

- Full keyboard navigation
- Visible focus indicators
- Screen-reader labels
- Semantic HTML
- Adjustable text size
- High-contrast mode
- Reduced-motion mode
- Ability to disable typing animation
- Captions or text equivalents for all audio
- No information conveyed only by color
- Touch targets large enough for mobile use
- A readable dialogue history
- Pause control for animated text
- Plain-language option for educational explanations where practical

The game should not require rapid responses.

19. Clinical and Ethical Safeguards

19.1 Outcome framing

Every final debrief should state that:

- Psychiatric outcomes are influenced by many factors.
- The simulation emphasizes one modifiable factor: clinician communication.
- A single phrase does not independently cause recovery, relapse, self-harm, hospitalization, or treatment discontinuation.
- The game compresses a complex clinical course for educational purposes.

19.2 Safety-related content

Cases involving suicide, self-harm, trauma, psychosis, coercion, or hospitalization should include:

- A content notice before the case
- Clinically appropriate language
- No graphic depictions
- No sensationalized endings
- A distinction between passive death wish, suicidal ideation, plan, intent, and behavior where relevant
- Clear explanation of safety assessment in the debrief
- Review by a qualified psychiatry educator before release

19.3 Patient portrayal

Patients should not function merely as diagnostic puzzles.

Each patient should have:

- Work, school, family, or community roles
- Preferences
- Values
- Strengths
- Coping strategies
- Sources of support
- Personal goals unrelated to symptom elimination
- Reactions that make sense from their perspective

Avoid defining characters solely by diagnosis.

20. Prototype Scope

Required for the first functional prototype

- Title screen
- Patient selection screen with three patients
- One complete, fully playable patient case
- Skeleton or abbreviated placeholder content for the other two
- Three encounters over six weeks
- Two static inter-encounter events
- Two binary trajectory resolutions
- Four complete endings for the functional case
- Hidden state tracking
- Encounter reflections
- Final debrief
- Local progress saving
- Responsive desktop and mobile interface
- Placeholder graphics
- Basic accessibility support
- Replay option
- Branch map after completion

Preferred for the next iteration

- All three cases fully written
- Twelve complete endings
- Educator mode
- Dialogue history
- Audio ambience
- Multiple portrait expressions
- Expanded accessibility settings
- Case completion dashboard
- Exportable learner reflection
- Instructor discussion guide

Out of scope for the initial build

- Generative AI dialogue
- Voice recognition
- Multiplayer
- Instructor accounts
- Learning-management-system integration
- Cloud saves
- User-created cases
- Analytics sent to an external server
- Automated clinical grading

- Procedurally generated patients
-

21. Development Sequence for Codex

Stage 1: Build the narrative engine

Create:

- Scene renderer
- Choice renderer
- State manager
- Effect processor
- Conditional scene logic
- Encounter resolver
- Ending resolver
- Local save system

Use simple placeholder text to verify that branching works.

Stage 2: Build the timeline system

Add:

- Week transitions
- Encounter transitions
- Life-event montages
- First and second branch storage
- Four-ending matrix

Confirm that every combination reaches the correct final state.

Stage 3: Build the interface

Add:

- Visual-novel conversation screen
- Patient portrait
- Background
- Dialogue box
- Response buttons
- Menu
- Mobile layout
- Accessibility controls

Stage 4: Add one complete demonstration case

The demonstration case should include:

- Fully authored Encounter 1
- Two variants of the first life event
- Two Encounter 2 openings
- Encounter 2 repair opportunities
- Four variants of the second life event
- Four Encounter 3 openings
- Four endings
- Encounter and final debriefs

Stage 5: Add patient selection and case templates

Create the three patient cards and provide structured placeholder data for the remaining cases.

Stage 6: Add replay and teaching features

Add:

- Completed-ending tracking
- Branch map
- Educator mode architecture
- Pivotal-choice review
- Alternative wording examples

Stage 7: Testing

Test:

- Every choice link
 - Every conditional scene
 - All four ending combinations
 - Save and resume
 - Restart behavior
 - Mobile layout
 - Keyboard navigation
 - Screen-reader labels
 - Missing assets
 - Invalid state recovery
 - Cases in which required safety flags are absent
-

22. Acceptance Criteria

The prototype is successful when:

1. The user can choose one of three patients.
 2. The user can complete one patient's six-week story from beginning to end.
 3. The story contains exactly three clinical encounters.
 4. Two fixed life events occur between the encounters.
 5. The player's dialogue choices alter the patient's responses to those events.
 6. Encounter 1 resolves into one of two hidden trajectories.
 7. Encounter 2 resolves into one of two hidden trajectories.
 8. The combination produces exactly four principal endings.
 9. A poor first encounter can be meaningfully repaired.
 10. A strong first encounter can still be undermined later.
 11. The player is not shown a visible relationship or morality score.
 12. Patient reactions provide immediate implicit feedback.
 13. Explicit educational feedback is provided after encounters and at the ending.
 14. The game saves progress locally.
 15. The game works on desktop and mobile.
 16. The content can be replaced or expanded without rewriting the engine.
 17. No backend or generative AI service is required.
 18. The game clearly frames patient outcomes as multifactorial.
 19. Placeholder art can later be replaced without changing case logic.
 20. All major pathways are reachable and testable.
-

23. Design Principles to Preserve

Codex should preserve these principles throughout implementation:

- **Consequences over quizzes:** The player learns by seeing what happens later.
 - **Patterns over single answers:** Branches emerge from cumulative communication.
 - **Repair over perfection:** A clinician can recognize and repair a rupture.
 - **Ambiguity over caricature:** Most choices should be plausible.
 - **Patient perspective over scoring:** The patient's response is the primary feedback.
 - **Agency over control:** The clinician influences but does not control the patient.
 - **Clinical realism over melodrama:** Outcomes should be consequential but credible.
 - **Data-driven architecture over hard-coded cases:** New patients should be easy to add.
 - **Reflection over punishment:** Even difficult endings should teach rather than shame.
 - **Originality over imitation:** The game may evoke the atmosphere of a conversation-focused visual novel without copying another game's assets, characters, writing, or interface.
-

24. Instruction to Codex

Build PsychiatryChat as a modular, data-driven, static web game.

Begin by implementing the complete engine and one demonstration patient with placeholder clinical content. Prioritize the six-week timeline, cumulative dialogue state, two-stage branching, static life events with dynamic reactions, four-ending structure, teaching debriefs, local saving, accessibility, and responsive interface.

Use clear content placeholders where final patient dialogue or artwork has not yet been supplied. Keep all clinically substantive text easy to locate and replace. Do not bury patient dialogue inside rendering functions.

The first prototype should demonstrate that the central mechanic works:

What the clinician says during one encounter changes what the patient discloses, believes, attempts, or avoids after leaving the room—and those consequences return to the clinical relationship weeks later.